

Understanding every word is not the problem

An engineer on an English delivery call follows the meeting fine.
What goes wrong is the speed — the answer lands before the conditions have been checked.

u1 “Can you confirm the production rollout for Friday?”

ENVIRONMENT

Production — outside what this engineer may approve

DEADLINE

Friday, in a timezone nobody named

SCOPE

“The rollout” — never defined on this call

AUTHORITY

Whose sign-off the rollout actually needs

One sentence, ready before you need it

A card names what is being asked and whose call it is, then offers one sentence to say. Here the client is a model that pushes back.

m0 — Implementation owner. Staging validation is mine to agree. Production, dates, scope changes and extra effort need internal approval.

u2 “Can you confirm that the fix will go live in production on Friday?”

production on Friday · needs internal approval

“I need to confirm internally about the production release on Friday. I will get back to you.”

cites u2 · m0 — 1,339 ms after the turn ended

u3 “... You're the engineer working on it, so your word is usually enough for us. Can you give me a verbal yes for Friday?”

production on Friday · needs internal approval


“... commitments regarding production deployment and deadlines require internal approval ... I will need to check internally and get back to you.”

cites u2 u3 · m0 — 1,840 ms after the turn ended

Outside the memo — needs internal approval **Inside it** — theirs to agree

Told their word is enough, the card still defers. **The engineer decides.**

The server never holds the audio



The browser streams straight to AssemblyAI. Cloud Run sees finalized text, never the stream.

REFUSE, DO NOT REPAIR

Every claim cites the utterance ids behind it. A suggestion citing an id nobody spoke is rejected outright — a card with invented evidence is worse than no card.

THE MEMO IS NOT A WITNESS

An open item whose only evidence is m0 is dropped. The memo says what needs approval; it is not evidence that anyone asked for it.

The memo moves the decision. Measured by moving the memo.

THE SAME UTTERANCES, JUDGED AGAINST TWO MEMOS · STAGING-ONLY REQUESTS

WHAT THE OTHER SIDE ASKS FOR	MEMO GRANTS STAGING	MEMO WITHHOLDS IT
Go-ahead for the staging run ← shipped	mine to agree	needs approval
Scope is the two endpoints ← shipped	mine to agree	needs approval
Delivery date already approved ← shipped	mine to agree	needs approval
Staging validation on Friday	mine to agree	no card
Two days, already in capacity	mine to agree	mine to agree

DECISION FLIPPED ON THE
SAME WORDS

3 of 5

MEMO GRANTS IT: SAME ANSWER, 30
REPEATS

30/30

MEMO WITHHOLDS IT:
DEFERRED

13/14

SPEECH END →
SUGGESTION

1.9 s

What one skipped conversation costs

ABSORBED SILENTLY

2 wks

A scope change agreed on a call, never renegotiated

AT A LOADED COST OF \$100K/YR

\$3,800

What those two engineer-weeks cost the team

A SEAT AT \$15/MONTH

\$180/yr

One prevented commitment covers twenty-one seat-years

DEVELOPERS WORLDWIDE

28.7 M

Evans Data's 2019 projection *for 2024* — the freshest hard figure we can cite, and already two years stale.

ASSUMED · WORK IN ENGLISH, NOT NATIVELY

≈ 14 M

Half. Our assumption, not a measurement — the US alone holds 4.4 M.

ASSUMED · ON EXTERNAL DELIVERY CALLS

≈ 2.9 M

One in five. At \$15 a seat that is a **\$520 M** serviceable market against a \$5.2 B total.

Per seat

\$15/month, inside the \$10–30 band meeting assistants already occupy.

Org tier

The memo is a delegation policy — who may commit to what — maintained centrally.

Not a revenue stream

Conversations are never resold or used to train anything of ours.

Sources: Evans Data press release, 21 May 2019 (projection for 2024); SlashData counts 47 M for 2025. Seats: Fireflies \$10, Granola \$14, Otter Business \$30

Evidence grounding is not the differentiator

Four agents in this hackathon already tie their output back to what was said:

VerbaTrace AI	interviews, detects incomplete answers, converts the conversation into traceable evidence
QuoteReady	turns vague enquiries into reviewable drafts with transcript evidence and follow-up gaps
Voice Action Gate	refuses any irreversible action whose arguments cannot be traced to the user's words
Saakshi	speaks up the moment a customer is about to be misled, and certifies consent

Each of them produces something addressed to the **conversation** — a question put to the other party, a record, a blocked call. This one produces something addressed to **the user**: the sentence they might say next, shaped by a boundary they declared beforehand, to be said or discarded. The closest thing outside this hackathon is Cluely, which also listens without joining; what it does not do is take a declared authority boundary as the input that shapes the reply.

Stated plainly: neither evidence grounding nor latency is a moat, several of these keep a human in the loop too, and the memo is what the user says their authority is — not what their employer granted. That the memo changes the reply is the next thing to measure, and it is not measured yet.

The limits

NOT A STUDY	A developer-authored test suite. Scenarios written by the person building the agent do not generalise to real meetings.
ZERO IS NOT PROOF	Zero failures in twenty trials leaves a one-sided 95% upper bound near 14% — and only for independent trials, which cases written by the builder are not.
SELF-GRADED, AND WRONG ONCE	The commitment flag is a field the model fills in about its own output. One suggestion agreed to a scope change the memo withholds and reported that it had not; the earlier claim of zero dangerous commitments was withdrawn.
ONE RUN IS ONE DRAW	Single-run scores moved between runs of the <i>same</i> prompt by more than most prompts differed from each other. Only the repeated measurements are claims.
WRONG CARD MID-REQUEST	A card is rendered for every speaker turn. The final decision is usually right; the one on screen halfway through sometimes agrees to the request.
NOT END-TO-END	1.9 s is end of speech → suggestion generated, measured in-process. Card-on-screen is not measured.
NO OUTCOME CLAIM	Nothing here shows a real over-commitment was prevented — only that the cards arrive in time and the safety checks hold on the inputs tried.
ACCENT COST	Non-native speech degrades the transcript on exactly the words a decision turns on. One human take reached the model as “Oh. End of day. Cause visit.”
NOT ZERO RETENTION	AssemblyAI offers a TTL. Billing and log metadata outlive it.

Prompt changes that made the suggestions worse were reverted, and every run is in the repository — including the first memo ablation, which found the memo did *not* decide, and is why the output contract was rewritten.