

Understanding every word is not the problem

An engineer on an English delivery call follows the meeting fine.
What goes wrong is the speed — the answer lands before the conditions have been checked.

u1 “Can you confirm the production rollout for Friday?”

ENVIRONMENT

Production — outside what this engineer may approve

DEADLINE

Friday, in a timezone nobody named

SCOPE

“The rollout” — never defined on this call

AUTHORITY

Whose sign-off the rollout actually needs

A reply, and whose decision it is

After each turn of the other side, a card names what is asked and whose call it is, then suggests a reply. Simulated client, TTS voices.

m0 — Implementation owner. Staging validation is mine to agree. Production, dates, scope changes and extra effort need internal approval.

u2 “Can you confirm that the fix will go live in production on Friday?”

production on Friday · **needs internal approval**


“I need to confirm internally about the production release on Friday. I will get back to you.”

cites u2 · m0 — suggestion call 1,339 ms

Outside the memo — needs internal approval **Inside it** — theirs to agree

The engineer reads it and decides. Nothing is spoken or sent.

Our server never receives the audio



The browser streams straight to AssemblyAI. Cloud Run sees finalized text, never the stream.

REFUSE, DO NOT REPAIR

Every claim cites the utterance ids behind it. A suggestion citing an id nobody spoke is rejected outright. The check is that the id exists, not that it supports the reply.

THE MEMO IS NOT A WITNESS

An open item whose only evidence is m0 is dropped. The memo says what needs approval; it is not evidence that anyone asked for it.

The memo moves the decision. Measured by moving the memo.

THE SAME UTTERANCES, JUDGED AGAINST TWO MEMOS · STAGING-ONLY REQUESTS

WHAT THE OTHER SIDE ASKS FOR		MEMO GRANTS STAGING	MEMO WITHHOLDS IT
Go-ahead for the staging run ← shipped		mine to agree	needs approval
Scope is the two endpoints ← shipped		mine to agree	needs approval
Delivery date already approved ← shipped		mine to agree	needs approval
Staging validation on Friday		mine to agree	no card
Two days, already in capacity		mine to agree	mine to agree
FLIPPED, SAME WORDS · 1 TEXT RUN · 1 NO CARD, 1 UNMOVED	MEMO GRANTS · LABELLED “MINE” · LABEL, NOT REPLY	MEMO WITHHOLDS · LABELLED “NEEDS APPROVAL”	AUDIO RUN · FINAL CARDS THAT AGREE TO WHAT THEY SHOULD NOT
3 of 5	15/15	13/15	4 of 20

WHO IT IS FOR • A HYPOTHESIS, NOT YET TESTED WITH BUYERS

The engineer feels it. The delivery lead pays for it.

USES IT

A Japanese-speaking engineer on English client calls, who follows the request but cannot approve dates, production or scope.

PAYS FOR IT

The delivery lead at a software services firm, who absorbs the unpaid rework when a quick yes becomes a commitment.

HAS TO ALLOW IT

Information security: client audio goes to AssemblyAI, the text and the memo to Gemini.

Price hypothesis: \$15 per seat per month. Willingness to pay, usage cost and avoided rework have not been measured.

What the engineer could use instead

INSTEAD	WHAT IT GIVES	WHAT IT MISSES
“Let me check and get back to you”	Safe, instant, free	Defers what the engineer may agree to as well
The memo, kept beside the call	The boundary in view	Matching it to the request, in English, mid-call
Copilot in Teams · Ask Gemini in Meet	Private answers about the meeting	Waits to be asked; Copilot only in meetings your own organisation hosts
This agent	A reply shaped by the boundary the engineer declared, after every turn of the other side	

Not yet measured: whether engineers prefer it to any of the three.

Three limits, and the check each needs next

WRONG AGREEMENTS

In the audio run, 4 of 20 final cards agree to what the memo withholds or never defined, and the model's own commitment flag has been wrong once. Next: score every card, not only the last.

OLD QUESTION

When the client changes the question after a refusal, the card can keep deferring on the old one. Next: test a fix against a model client that argues back.

NOT TRIED LIVE

The cases were written by the builder, and no engineer has read a card during a live call. Next: mock calls against a memorised line and a memo on paper.

Latency, accent errors, retention and reverted prompts: docs/eval in the repository.