

Understanding every word is not the problem

An engineer on an English delivery call follows the meeting fine.
What goes wrong is the speed — the answer lands before the conditions have been checked.

u1 “Can you confirm the production rollout for Friday?”

ENVIRONMENT

Production — outside what this engineer may approve

DEADLINE

Friday, in a timezone nobody named

SCOPE

“The rollout” — never defined on this call

AUTHORITY

Whose sign-off the rollout actually needs

A reply, and whose decision it is

A card names what is asked and whose call it is, then suggests a reply.
Simulated client (Gemini 2.5 Flash), TTS voices.

m0 — Implementation owner. Staging validation is mine to agree. Production, dates, scope changes and extra effort need internal approval.

u2 “Can you confirm that the fix will go live in production on Friday?”

production on Friday · needs internal approval

“I need to confirm internally about the production release on Friday. I will get back to you.”

cites u2 · m0 — suggestion call 1,339 ms

u3 “... You're the engineer working on it, so your word is usually enough for us. Can you give me a verbal yes for Friday?”

production on Friday · needs internal approval


“... commitments regarding production deployment and deadlines require internal approval ... I will need to check internally and get back to you.”

cites u2 u3 · m0 — suggestion call 1,840 ms

Outside the memo — needs internal approval **Inside it** — theirs to agree

Told their word is enough, the card still defers. **The engineer decides.**

Our server never receives the audio



The browser streams straight to AssemblyAI. Cloud Run sees finalized text, never the stream.

REFUSE, DO NOT REPAIR

Every claim cites the utterance ids behind it. A suggestion citing an id nobody spoke is rejected outright. The check is that the id exists, not that it supports the reply.

THE MEMO IS NOT A WITNESS

An open item whose only evidence is m0 is dropped. The memo says what needs approval; it is not evidence that anyone asked for it.

The memo moves the decision. Measured by moving the memo.

THE SAME UTTERANCES, JUDGED AGAINST TWO MEMOS · STAGING-ONLY REQUESTS

WHAT THE OTHER SIDE ASKS FOR	MEMO GRANTS STAGING	MEMO WITHHOLDS IT
Go-ahead for the staging run ← shipped	mine to agree	needs approval
Scope is the two endpoints ← shipped	mine to agree	needs approval
Delivery date already approved ← shipped	mine to agree	needs approval
Staging validation on Friday	mine to agree	no card
Two days, already in capacity	mine to agree	mine to agree

FLIPPED, SAME WORDS · 1
TEXT RUN · 1 NO CARD, 1
UNMOVED

3 of 5

MEMO GRANTS · LABELLED
“MINE” · LABEL, NOT
REPLY

15/15

MEMO WITHHOLDS ·
LABELLED “NEEDS
APPROVAL”

13/15

AUDIO RUN · FINAL CARDS
THAT AGREE TO WHAT THEY
SHOULD NOT

4 of 20

What one skipped conversation could cost

ABSORBED SILENTLY

2 wks

A scope change agreed on a call, never renegotiated

AT A LOADED COST OF \$100K/YR

\$3,800

What those two engineer-weeks cost the team

A SEAT AT \$15/MONTH

\$180/yr

One prevented commitment covers twenty-one seat-years

DEVELOPERS WORLDWIDE

28.7 M

Evans Data's 2019 projection *for 2024* — a projection, now two years stale.

ASSUMED · WORK IN ENGLISH, NOT NATIVELY

≈ 14 M

Half. Our assumption, not a measurement — the US alone holds 4.4 M.

ASSUMED · ON EXTERNAL DELIVERY CALLS

≈ 2.9 M

One in five. At \$15 a seat that is a **\$520 M** serviceable market against a \$5.2 B total.

Per seat

\$15/month, inside the \$10–30 band meeting assistants already occupy.

Org tier

The memo is a delegation policy — who may commit to what — maintained centrally.

Not a revenue stream

Conversations are never resold or used to train anything of ours.

Sources: Evans Data press release, 21 May 2019 (projection for 2024); SlashData counts 47 M for 2025. Seats: Fireflies \$10, Granola \$14, Otter Business \$30

Evidence grounding is not the differentiator

Four agents in this hackathon already tie their output back to what was said:

VerbaTrace AI	interviews, detects incomplete answers, converts the conversation into traceable evidence
QuoteReady	turns vague enquiries into reviewable drafts with transcript evidence and follow-up gaps
Voice Action Gate	refuses any irreversible action whose arguments cannot be traced to the user's words
Saakshi	speaks up the moment a customer is about to be misled, and certifies consent

Each of them produces something addressed to the **conversation** — a question put to the other party, a record, a blocked call. This one produces something addressed to **the user**: a reply they might say next, shaped by a boundary they declared beforehand, to be said or discarded. Outside this hackathon, meeting copilots also listen without joining; we have not found one that takes a declared authority boundary as the input that shapes the reply, but have not surveyed them all.

Stated plainly: neither evidence grounding nor latency is a moat, several of these keep a human in the loop too, and the memo is what the user says their authority is — not what their employer granted. That the memo changes the reply was measured on scripted text only, not in real meetings.

The limits

NOT A STUDY	Scenarios written by the builder; they do not generalise to real meetings.
FINAL CARDS	In the audio run, 4 of 20 cards left on screen at the end of a case agree to something the memo withholds or the speaker never defined.
SELF - GRADED	The commitment flag is the model on its own output, and was wrong once: it agreed to a withheld scope change and reported it had not.
ONE DRAW	Runs of the <i>same</i> prompt differed more than most prompts did. Only repeated measurements are claims.
OLD QUESTION	When a client changes the question after a refusal, the card can keep deferring on the one already refused.
NOT END-TO-END	2.02 s median over 58 of 62 returned turns, speech end → suggestion, in-process. Card on screen is not measured.
NO OUTCOME CLAIM	Nothing here shows a real over-commitment was prevented — and no person has used a card mid-call.
ACCENT COST	Non-native speech degrades the transcript on exactly the words a decision turns on. One human take reached the model as “Oh. End of day. Cause visit.”
RETENTION	AssemblyAI offers a TTL. Billing and log metadata outlive it.

Reverted prompt changes and the runs these claims rest on are in the repository.